

# SGD for linear regression

February 19, 2019

## 1 Assignment 6: Implement SGD for linear regression

### The Boston Housing Dataset

A Dataset derived from information collected by the U.S. Census Service concerning housing in the area of Boston Mass.

*This dataset contains information collected by the U.S Census Service concerning housing in the area of Boston Mass. It was obtained from the StatLib archive (<http://lib.stat.cmu.edu/datasets/boston>), and has been used extensively throughout the literature to benchmark algorithms. However, these comparisons were primarily done outside of Delve and are thus somewhat suspect. **The dataset is small in size with only 506 cases***

**Dataset Naming** The name for this dataset is simply boston. It has two prototasks: nox, in which the nitrous oxide level is to be predicted; and price, in which the median value of a home is to be predicted

### Dataset Information

- **Origin** The origin of the boston housing data is Natural.
- **Usage** This dataset may be used for Assessment.
- **Number of Cases** The dataset contains a total of 506 cases.
- **Order** The order of the cases is mysterious.
- **Variables** There are 14 attributes in each case of the dataset. They are: CRIM - per capita crime rate by town ZN - proportion of residential land zoned for lots over 25,000 sq.ft. INDUS - proportion of non-retail business acres per town. CHAS - Charles River dummy variable (1 if tract bounds river; 0 otherwise) NOX - nitric oxides concentration (parts per 10 million) RM - average number of rooms per dwelling AGE - proportion of owner-occupied units built prior to 1940 DIS - weighted distances to five Boston employment centres RAD - index of accessibility to radial highways TAX - full-value property-tax rate per 10,000dollars PTRATIO - pupil-teacher ratio by town B -  $1000(B_k - 0.63)^2$  where  $B_k$  is the proportion of blacks by town LSTAT - percent lower status of the population MEDV - Median value of owner-occupied homes in \$1000's

### Importing the necessary libraries

```
In [50]: import warnings
          warnings.filterwarnings("ignore", category=DeprecationWarning)

          import pandas as pd
          import numpy as np
```

```

from sklearn.datasets import load_boston
from sklearn.preprocessing import StandardScaler
from sklearn.linear_model import SGDRegressor
from matplotlib import pyplot as plt
import seaborn as sns
from sklearn.metrics import mean_squared_error
from tqdm import tqdm
import random
from sklearn.model_selection import train_test_split

```

### *Loading the dataset*

```

In [51]: boston = load_boston()
        X = pd.DataFrame(boston.data)
        Y = np.array(boston.target)

        output_df = pd.DataFrame(columns=['Algorithm', 'Learning Rate', 'N-Iteration', 'M.S.E'])
        n_iter_range = [10**i for i in range(0, 6)]

```

### *Splitting the dataset with 70-30 train-test ratio*

```

In [52]: XTrain, XTest, y_train, y_test = train_test_split(X, Y, test_size=0.30, random_state=0,

```

### *Standardizing the dataset*

```

In [53]: # Standardizing the dataset with mean centering and variance scaling
        sc = StandardScaler()
        #Fitting and transforming the training dataset
        x_train = sc.fit_transform(XTrain)
        #Transforming the testing dataset
        x_test = sc.transform(XTest)

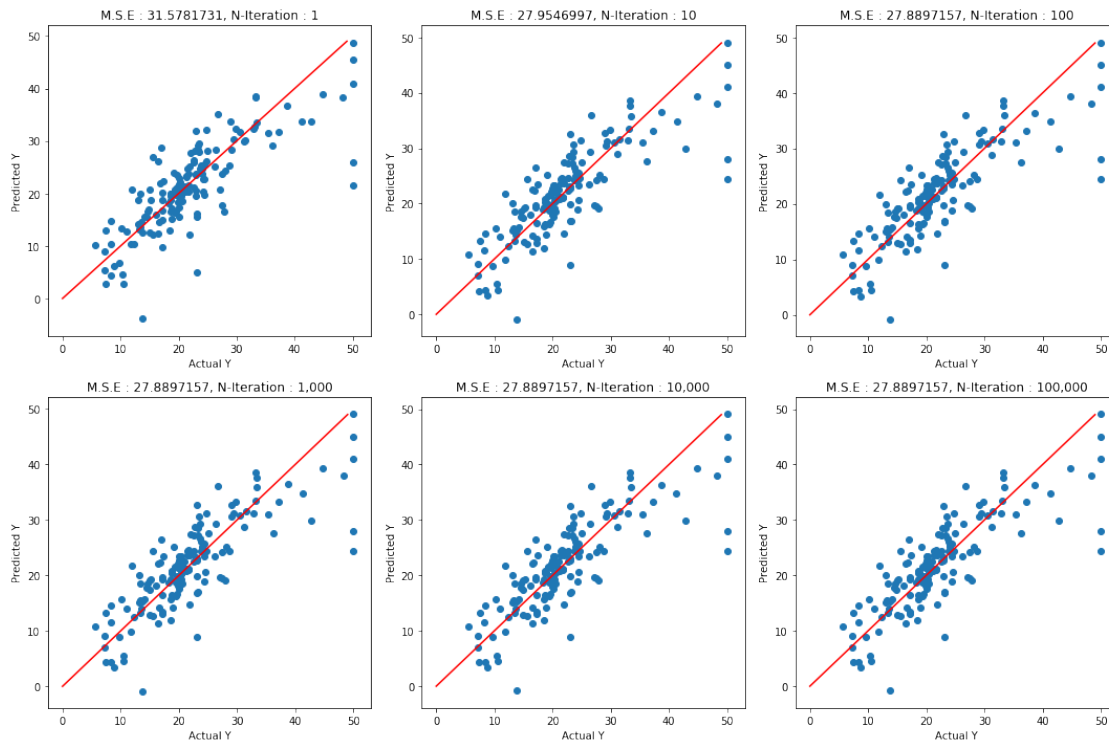
```

## **1.1 SGD Regressor with Learning Rate : 0.01**

```


In [54]: f, axarr = plt.subplots(2, 3, figsize=(15,10))
        save_sgd_df = pd.DataFrame(columns=['Algorithm', 'Learning Rate', 'N-Iteration', 'M.S.E'])
        for index, ite in enumerate(n_iter_range):
            clf = SGDRegressor(penalty='none', learning_rate='constant', eta0=0.01, shuffle=False)
            clf.fit(x_train, y_train)
            y_pred = clf.predict(x_test)
            mse = mean_squared_error(y_test, y_pred)
            save_sgd_df = save_sgd_df.append({'Algorithm': 'SGDRegressor', 'Learning Rate': 0.01, 'N-Iteration': ite, 'M.S.E': mse})
            axarr[int(index/3), int(index%3)].plot(np.arange(50), np.arange(50), color='r')
            axarr[int(index/3), int(index%3)].scatter(y_test, y_pred)
            axarr[int(index/3), int(index%3)].set_xlabel('Actual Y')
            axarr[int(index/3), int(index%3)].set_ylabel('Predicted Y')
            axarr[int(index/3), int(index%3)].set_title('M.S.E : {0:.7f}'.format(mse) + ', N-Iteration: ' + str(ite))
        plt.tight_layout()
        plt.grid()
        plt.show()

```



*Plotting the N-Iterations vs M.S.E for SGDRegressor Algorithm*

```
In [55]: plt.figure(figsize=(18,2))
plt.scatter(save_sgd_df['N-Iteration'], save_sgd_df['M.S.E'])
plt.plot(save_sgd_df['N-Iteration'], save_sgd_df['M.S.E'], '-o')
plt.grid()
plt.show()
```



*Printing the Weight Vector for SGD Regressor*

```
In [56]: sgd_op_w = pd.DataFrame(clf.coef_).T
sgd_op_w.index = ['SGDRegressor']
sgd_op_w
```

```
Out[56]:
```

	0	1	2	3	4	5
SGDRegressor	-1.191457	1.241887	0.269344	1.286084	-1.871315	3.039862

```

        6          7          8          9          10          11 \
SGDRegressor -0.139336 -3.418594  2.470299 -1.982521 -2.657993  0.787101

        12
SGDRegressor -3.644444

```

## 1.2 Implementing SDG Regressor with learning rate : 0.01 from scratch

```

In [57]: plt.close()
         f, axarr = plt.subplots(2, 3, figsize=(15,10))

         learning_rate = 0.01
         N = x_train.shape[0]
         features = x_train.shape[1]

         save_mnlsgd_df = pd.DataFrame(columns=['Algorithm', 'Learning Rate', 'N-Iteration', 'M.S.E'])
         #Selecting the batch size for SGD to be constant as 50% of training dataset
         k = int(0.5 * N)

         for index,ite in enumerate(n_iter_range):
             #Randomly selecting the gradient_w vector from the data itself
             w_optimal = np.zeros(shape=(features,))
             #Selecting Intercept term as 0
             b_optimal = 0

             for it in np.arange(ite):
                 res_w = np.zeros((features,))
                 res_b = 0
                 batch_li = np.array([random.randint(1, N-1) for _ in range(k)])
                 for i in batch_li:
                     y_pred = np.dot(w_optimal, x_train[i]) + b_optimal
                     err = y_train[i] - y_pred
                     res_w = np.array(res_w + (err * x_train[i]), dtype=np.float64)
                     res_b = res_b + err
                 grad_w = (-2 / k) * res_w
                 grad_b = (-2 / k) * res_b
                 w_optimal = w_optimal - (learning_rate * grad_w)
                 b_optimal = b_optimal - (learning_rate * grad_b)

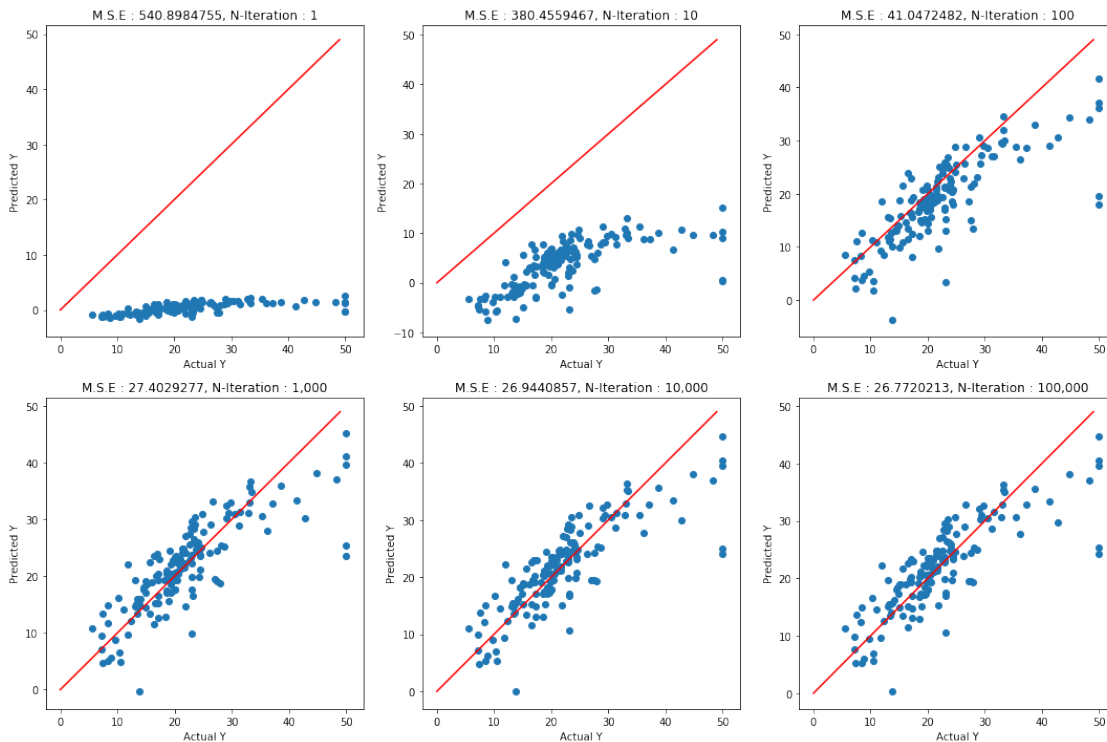
             y_pred_ = []
             for i in range(len(x_test)):
                 y=np.asscalar(np.dot(w_optimal,x_test[i]) + b_optimal)
                 y_pred_.append(y)
             mse = mean_squared_error(y_test, y_pred_)
             save_mnlsgd_df = save_mnlsgd_df.append({'Algorithm': 'Manual SGD', 'Learning Rate': learning_rate, 'N-Iteration': it, 'M.S.E': mse})
             axarr[int(index/3), int(index%3)].plot(np.arange(50), np.arange(50), color='r')
             axarr[int(index/3), int(index%3)].scatter(y_test, y_pred_)
             axarr[int(index/3), int(index%3)].set_xlabel ('Actual Y')

```

```

axarr[int(index/3), int(index%3)].set_ylabel ('Predicted Y')
axarr[int(index/3), int(index%3)].set_title('M.S.E : {0:.7f}'.format(mse) + ' , N-I')
plt.tight_layout()
plt.grid()
plt.show()

```




*Plotting the N-Iterations vs M.S.E for ManualSGD Algorithm*

```

In [58]: plt.figure(figsize=(18,2))
plt.scatter(save_mnlskd_df['N-Iteration'], save_mnlskd_df['M.S.E'])
plt.plot(save_mnlskd_df['N-Iteration'], save_mnlskd_df['M.S.E'], '-o')
plt.grid()
plt.show()

```



*Printing the Weight Vector for Manual/Scratch SGD Regressor*

```

In [59]: man_sgd_op_w = pd.DataFrame(w_optimal).T
man_sgd_op_w.index = ['Manual-SGD']
man_sgd_op_w

```

```
Out [59]:
```

	0	1	2	3	4	5	\
Manual-SGD	-0.955591	1.010319	0.080567	0.642676	-1.800255	2.662878	
	6	7	8	9	10	11	\
Manual-SGD	-0.224925	-3.115612	2.213754	-1.886661	-2.261027	0.612757	
	12						
Manual-SGD	-3.651006						

```
In [60]: out_resp = pd.DataFrame(columns=['Algorithm', 'Learning Rate', 'N-Iteration', 'M.S.E'])
for i in range(6):
    out_resp = out_resp.append(save_sgd_df.iloc[i])
    out_resp = out_resp.append(save_mnlsgd_df.iloc[i])

#Reporting the weight vector result
w_out_resp = pd.concat([sgd_op_w, man_sgd_op_w])
```

### 1.3 Results

#### *Checking the Weight Vector for both algorithm implementation*

```
In [61]: w_out_resp
```

```
Out [61]:
```

	0	1	2	3	4	5	\
SGDRegressor	-1.191457	1.241887	0.269344	1.286084	-1.871315	3.039862	
Manual-SGD	-0.955591	1.010319	0.080567	0.642676	-1.800255	2.662878	
	6	7	8	9	10	11	\
SGDRegressor	-0.139336	-3.418594	2.470299	-1.982521	-2.657993	0.787101	
Manual-SGD	-0.224925	-3.115612	2.213754	-1.886661	-2.261027	0.612757	
	12						
SGDRegressor	-3.644444						
Manual-SGD	-3.651006						

#### *Checking the MeanSquaredError (M.S.E) for both algorithm implementations for all N-Iterations*

```
In [62]: out_resp
```

```
Out [62]:
```

	Algorithm	Learning Rate	N-Iteration	M.S.E
0	SGDRegressor	0.01	1	31.578173
0	Manual SGD	0.01	1	540.898476
1	SGDRegressor	0.01	10	27.954700
1	Manual SGD	0.01	10	380.455947
2	SGDRegressor	0.01	100	27.889716
2	Manual SGD	0.01	100	41.047248
3	SGDRegressor	0.01	1000	27.889716
3	Manual SGD	0.01	1000	27.402928

4	SGDRegressor	0.01	10000	27.889716
4	Manual SGD	0.01	10000	26.944086
5	SGDRegressor	0.01	100000	27.889716
5	Manual SGD	0.01	100000	26.772021

## 1.4 Conclusion

- As the number of iterations for the manual/scratch SGD regressor algorithm is increasing, the Mean Squared Error (M.S.E) decreases.
- At nearly 1k iterations, our manual/scratch SGD Regressor algorithm MSE value becomes significantly equal to the SGDRegressor algorithm.
- We can see that our manual/scratch SGD Regressor algorithm performs well and is implemented correctly.