

Assignment 2 - TSNE Amazon Fine Food Reviews

November 28, 2018

1 Amazon Fine Food Reviews

Context

This dataset consists of reviews of fine foods from amazon. The data span a period of more than

Dataset Description

The dataset is included with 2 formats - Reviews.csv and database.sqlite(which contains Reviews

Overall Description:

Data includes:

- Reviews from Oct 1999 - Oct 2012
- 568,454 reviews
- 256,059 users
- 74,258 products
- 260 users with > 50 reviews

High level Description: Attributes Information

- Id - Row Id
- ProductId - Unique identifier for the product
- UserId - Unique identifier for the user
- ProfileName - Profile name of the user
- HelpfulnessNumerator - Number of users who found the review helpful
- HelpfulnessDenominator - Number of users who indicated whether they found the review helpful
- Score - Rating between 1 and 5
- Time - Timestamp for the review
- Summary - Brief summary of the review
- Text - Text of the review

Objective

The main objective is - Given a review, determine whether that review is Positive or Negative.

The review can be classified as positive if the Rating/Score is more than 3 and similarly, negative if it is less than 3.

Note - This is a very approximate and proxy method of determining the polarity of the reviews as

1.1 Importing the libraries

```
In [1]: %matplotlib inline
import warnings
warnings.filterwarnings('ignore')

import numpy as np
import pandas as pd
import seaborn as sns
from matplotlib import pyplot as plt
from tqdm import tqdm
import string
import sqlite3
import nltk
import re
import time
from nltk.corpus import stopwords

from sklearn.preprocessing import StandardScaler
from sklearn.manifold import TSNE
from sklearn.feature_extraction.text import CountVectorizer
from sklearn.feature_extraction.text import TfidfVectorizer

from gensim.models import Word2Vec
```

1.2 Loading the data

```
In [2]: # Create a connection
conn = sqlite3.connect('../..../Datasets/amazon-reviews-dataset/database.sqlite')

# Read from the sqlite3 database using the connection and save it in dataframe using pandas
# Note - We will only fetch those reviews where the score is not equal to 3 (Assumption)
data_df = pd.read_sql_query('SELECT * FROM REVIEWS WHERE SCORE != 3', conn)

# Printing out the shape of the dataframe
df_shape_tup = data_df.shape
print("Number of Reviews - {} and Features/Attributes - {}".format(df_shape_tup[0], df_shape_tup[1]))
```

Number of Reviews - 525814 and Features/Attributes - 10

```
In [3]: # Printing out the first 2 rows of data_db dataframe
data_df.head(2)
```

```
Out[3]:
```

	Id	...	
0	1	...	I have bought several of the V
1	2	...	Product arrived labeled as Jum

[2 rows x 10 columns]


```
In [4]: # Printing out the higher level statistics of the dataset
data_df.describe()
```

```
Out[4]:
```

	Id	...	Time
count	525814.000000	...	5.258140e+05
mean	284599.060038	...	1.295943e+09
std	163984.038077	...	4.828129e+07
min	1.000000	...	9.393408e+08
25%	142730.250000	...	1.270598e+09
50%	284989.500000	...	1.310861e+09
75%	426446.750000	...	1.332634e+09
max	568454.000000	...	1.351210e+09

```
[8 rows x 5 columns]
```

```
In [3]: plt.close()
sns.distplot(data_df['Score'].values)
plt.grid()
plt.show()
```



```
In [5]: # Printing out the higher level information of each feature
print(data_df.info())
print("\nThe Unique values in the class attribute column are : {}".format(data_df['Score'].unique()))
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 525814 entries, 0 to 525813
```

```
Data columns (total 10 columns):
Id                525814 non-null int64
ProductId         525814 non-null object
UserId           525814 non-null object
ProfileName       525814 non-null object
HelpfulnessNumerator  525814 non-null int64
HelpfulnessDenominator 525814 non-null int64
Score            525814 non-null int64
Time             525814 non-null int64
Summary          525814 non-null object
Text             525814 non-null object
dtypes: int64(5), object(5)
memory usage: 40.1+ MB
None
```

The Unique values in the class attribute column are : [5 1 4 2]

Observation: 1. The dataset contains 525814 rows with 10 attributes/features. 2. The dataset doesnot contain any null/NA values in features. 3. For our analysis, the score is used for classification whether the review 'Text' is positive or negative. Since the column/attribute/feature is Numeric. It must be converted to categorical datatype by using some if else logic. 4. The Score can be mapped as Positive if the value is greater than 3 and Negative if the value is less than 3.

1.3 Data Preprocessing, Data DeDuplication, Stemming, Removing Stopwords

```
In [6]: # Data Preprocessing
# 1. Score mapping - The Score which is classification attribute needs to be mapped to Positive or Negative
# for better classification visually.
```

```
def score_mapping(score):
    if score > 3:
        return "Positive"
    else:
        return "Negative"

data_df['Score'] = data_df['Score'].map(score_mapping)
print(data_df['Score'].value_counts())
print("\n" + str(data_df['Score'].value_counts(normalize=True)))
print("\n" + str(data_df.isnull().any()))
```

```
Positive    443777
Negative    82037
Name: Score, dtype: int64
```

```
Positive    0.843981
Negative    0.156019
```

Name: Score, dtype: float64

```
Id                False
ProductId         False
UserId           False
ProfileName       False
HelpfulnessNumerator  False
HelpfulnessDenominator False
Score            False
Time             False
Summary          False
Text             False
dtype: bool
```

Observation: 1. As per the analysis on class attribute "Score", we find that the dataset is quite imbalanced with 84% values as positive and 16% values as negative. 2. Also there are no null values in any of the columns/attributes/features.

```
In [7]: # Data Preprocessing
        # 2. Helpfulness Numerator must be less than or equal to Helpfulness Denominator
        previous_n = data_df.shape[0]
        data_df = data_df[data_df['HelpfulnessNumerator'] <= data_df["HelpfulnessDenominator"]]
        print("After Data processing Step 2: Shape - " + str(data_df.shape))
        print("{} reviews removed".format(previous_n-data_df.shape[0]))
```

After Data processing Step 2: Shape - (525812, 10)
2 reviews removed

```
In [8]: # Data Preprocessing
        # 3. Removing duplicate entries
        previous_n = data_df.shape[0]
        data_df.sort_values(by='ProductId', axis=0, ascending=True, inplace=True, kind='quicksort')
        data_df.drop_duplicates(subset={"UserId","ProfileName","Time","Text"}, keep='first', inplace=True)
        print("After Data processing Step 2: Shape - " + str(data_df.shape))
        print("{} reviews removed".format(previous_n-data_df.shape[0]))
```

After Data processing Step 2: Shape - (364171, 10)
161641 reviews removed

```
In [9]: # Data Preprocessing
        # 4. Removing HTML tags,punctuations,stopwords and stemming
        def CleanHTML(sentence):
            cleanedText = re.sub('<.*?>', '', sentence)
            return cleanedText

        def CleanPunc(sentence):
```

```

cleanedText = re.sub(r'[?!|@|#|^|&|*|_|-|"\|\'|:|;|+]', r'', sentence)
cleanedText = re.sub(r'[.,|)|(|\|/]', r' ', cleanedText)
return cleanedText

def cleanText(input):
    stop = set(stopwords.words('english')) #set of stopwords
    sno = nltk.stem.SnowballStemmer('english') #initialising the snowball stemmer
    cleanedText_df = pd.DataFrame()
    final_str = []
    str = ''
    for sentence in tqdm(input.values):
        filtered_sentence = []
        for words in CleanHTML(sentence).split():
            for cleanedText in CleanPunc(words).split():
                if(len(cleanedText) > 2 & cleanedText.isalpha()):
                    if(cleanedText.lower() not in stop):
                        s = (sno.stem(cleanedText.lower())).encode('UTF-8')
                        filtered_sentence.append(s)
        str = b' '.join(filtered_sentence)
        final_str.append(str)
    cleanedText_df['CleanedText'] = final_str
    cleanedText_df['CleanedText'] = cleanedText_df['CleanedText'].str.decode('UTF-8')
    return cleanedText_df['CleanedText'].values

scores_df = data_df['Score'].values
final_df = pd.DataFrame(data_df['Text'], columns=['Text'])
final_df['Score'] = scores_df

del data_df
del scores_df
final_df = final_df.sample(n=15000, random_state=0)
final_df['CleanedText'] = cleanText(final_df['Text'])

final_df.shape

```

100%| 15000/15000 [00:13<00:00, 1142.28it/s]

Out[9]: (15000, 3)

```

In [10]: # BOW
def bowUniGramMethod(df):
    print("Starting BOW methods with shape : " + str(df.shape))
    # BOW Uni-Gram
    scores_df = df['Score']
    start_time = time.time()
    print('Starting BoW UniGram')

```

```

final_text = CountVectorizer().fit_transform(df['CleanedText'].values)
print("The Shape of BoW Model Uni-Gram : " + str(final_text.shape))

print('Converting the csr_matrix to ndarray')
final_text = final_text.toarray().astype(float)
print('Standardizing the data')
final_text = StandardScaler().fit_transform(final_text)
print('Starting the TSNE Model')
tsne_data = TSNE(perplexity=30, n_iter=5000, n_components=2, random_state=0).fit_tr
final_df = pd.DataFrame(tsne_data, columns=['dim1', 'dim2'])
final_df['Score'] = scores_df.values
print('TSNE Model completed for Bow-UniGram in ' + str(time.time() - start_time))

plt.close()
sns.FacetGrid(data=final_df, hue='Score', height=10).map(plt.scatter, "dim1", "dim2")
plt.title('TSNE - Amazon Review Dataset with (15k points, 30 Perplexity, 5k iterati
plt.grid()
plt.show()


```

```
bowUniGramMethod(final_df)
```

```

Starting BOW methods with shape : (15000, 3)
Starting BoW UniGram
The Shape of BoW Model Uni-Gram : (15000, 19849)
Converting the csr_matrix to ndarray
Standardizing the data
Starting the TSNE Model
TSNE Model completed for Bow-UniGram in 10120.856751441956

```



```
In [11]: # TF-IDF
def tfidfUniGramMethod(df):
    print("Starting TF_IDF methods with shape : " + str(df.shape))
    global dictionary
    # TF_IDF Uni-Gram
    scores_df = df['Score']
    start_time = time.time()
    print('Starting TF_IDF UniGram')
    model = TfidfVectorizer()
    final_text = model.fit_transform(df['CleanedText'].values)
    print("The Shape of TF_IDF Model Uni-Gram : " + str(final_text.shape))
    dictionary = dict(zip(model.get_feature_names(), list(model.idf_)))

    print('Converting the csr_matrix to ndarray')
    final_text = final_text.toarray().astype(float)
    print('Standardizing the data')
```



```

final_text = StandardScaler().fit_transform(final_text)
print('Starting the TSNE Model')
tsne_data = TSNE(perplexity=30, n_iter=5000, n_components=2, random_state=0).fit_tr
final_df = pd.DataFrame(tsne_data, columns=['dim1', 'dim2'])
final_df['Score'] = scores_df.values
print('TSNE Model completed for TF_IDF-UniGram in '+ str(time.time() - start_time))

plt.close()
sns.FacetGrid(data=final_df, hue='Score', height=10).map(plt.scatter, "dim1", "dim2")
plt.title('TSNE - Amazon Review Dataset with (15k points, 30 Perplexity, 5k iterati
plt.grid()
plt.show()


```

```
tfidfUniGramMethod(final_df)
```

```

Starting TF_IDF methods with shape : (15000, 3)
Starting TF_IDF UniGram
The Shape of TF_IDF Model Uni-Gram : (15000, 19849)
Converting the csr_matrix to ndarray
Standardizing the data
Starting the TSNE Model
TSNE Model completed for TF_IDF-UniGram in 10463.602090597153

```



```
In [12]: def returnListOfSentences(df):
    list_of_sentences = []
    for sent in df.values:
        list_of_sentences.append(sent.split())
    return list_of_sentences

start_time = time.time()
sentList = returnListOfSentences(final_df['CleanedText'])
print("Starting W2V Model..")
w2v_model = Word2Vec(sentList, size=100, min_count=4, workers=8)
w2v_words = list(w2v_model.wv.vocab)
print("Number of words that occurred minimum 4 times ", len(w2v_words))
#print("sample words ", w2v_words[0:50])
print('W2V Model completed for in ' + str(time.time() - start_time))
```

Starting W2V Model..

Number of words that occurred minimum 4 times 6824

W2V Model completed for in 4.378483533859253

```
In [13]: def avgW2VcomputeModel(list_of_sent):
    # average Word2Vec
    # compute average word2vec for each review.
    start_time = time.time()
    sent_vectors = []; # the avg-w2v for each sentence/review is stored in this list
    for sent in tqdm(list_of_sent): # for each review/sentence
        sent_vec = np.zeros(100) # as word vectors are of zero length
        cnt_words = 0; # num of words with a valid vector in the sentence/review
        for word in sent: # for each word in a review/sentence
            if word in w2v_words:
                vec = w2v_model.wv[word]
                sent_vec += vec
                cnt_words += 1
        if cnt_words != 0:
            sent_vec /= cnt_words
        sent_vectors.append(sent_vec)
    print('AvgW2V Model completed for in ' + str(time.time() - start_time))
    return sent_vectors
```

```
print("Computing AvgWord2Vec Model..")
sentVectors = avgW2VcomputeModel(sentList)
```

```
1%|          | 142/15000 [00:00<00:10, 1404.58it/s]
```

Computing AvgWord2Vec Model..

```
100%|| 15000/15000 [00:22<00:00, 654.65it/s]
```

AvgW2V Model completed for in 22.920954942703247

```
In [14]: def plotTSNEModelForAvgW2V(sentVect, score_df):
    start_time = time.time()
    print("Converting to ndarray")
    arrayVec = np.array(sentVect, dtype=np.float64)
    print("Standardizing Array")
    standardizedVec = StandardScaler().fit_transform(arrayVec)
    print("Starting TSNE Model..")
    tsne_data = TSNE(perplexity=30, n_iter=5000, n_components=2, random_state=0).fit_tr
    final_df = pd.DataFrame(tsne_data, columns=['dim1', 'dim2'])
    final_df['Score'] = score_df.values
```

```

print('TSNE Model completed for AvgW2V in ' + str(time.time() - start_time))

plt.close()
sns.FacetGrid(data=final_df, hue='Score', height=10).map(plt.scatter, "dim1", "dim2")
plt.title('TSNE - Amazon Review Dataset with (15k points, 30 Perplexity, 5k iterations)')
plt.grid()
plt.show()

```


```
plotTSNEModelForAvgW2V(sentVectors, final_df['Score'])
```

Converting to ndarray

Standardizing Array

Starting TSNE Model..

TSNE Model completed for AvgW2V in 2007.5378654003143



```

In [15]: def tfidfW2VComputeModel(list_of_sent):
    # TF_IDF Word2Vec
    # compute TF_IDF word2vec for each review.
    start_time = time.time()

    tfidf_sent_vectors = []; # the tf-idf-w2v for each sentence/review is stored in the
    for sent in tqdm(list_of_sent): # for each review/sentence
        sent_vec = np.zeros(100) # as word vectors are of zero length
        weight_sum = 0; # num of words with a valid vector in the sentence/review
        for word in sent: # for each word in a review/sentence
            if word in w2v_words:
                if word in dictionary:
                    vec = w2v_model.wv[word]
                    tf_idf = dictionary[word]*(sent.count(word)/len(sent))
                    sent_vec += (vec * tf_idf)
                    weight_sum += tf_idf
        if weight_sum != 0:
            sent_vec /= weight_sum
        tfidf_sent_vectors.append(sent_vec)
    print('TF_IDFW2V Model completed for in ' + str(time.time() - start_time))
    return tfidf_sent_vectors

print("Computing TF_IDFW2V Model..")
sentVectors = tfidfW2VComputeModel(sentList)

1%|          | 107/15000 [00:00<00:14, 1049.89it/s]

Computing TF_IDFW2V Model..

100%|| 15000/15000 [00:36<00:00, 415.22it/s]

TF_IDFW2V Model completed for in 36.130290031433105

```

```

In [16]: def plotTSNEModelForTFIDFW2V(sentVect, score_df):
    start_time = time.time()
    print("Converting to ndarray")
    arrayVec = np.array(sentVect, dtype=np.float64)
    print("Standardizing Array")
    standardizedVec = StandardScaler().fit_transform(arrayVec)
    print("Starting TSNE Model..")
    tsne_data = TSNE(perplexity=30, n_iter=5000, n_components=2, random_state=0).fit_tr
    final_df = pd.DataFrame(tsne_data, columns=['dim1', 'dim2'])
    final_df['Score'] = score_df.values
    print('TSNE Model completed for TF-IDFW2V in ' + str(time.time() - start_time))

```

```
plt.close()
sns.FacetGrid(data=final_df, hue='Score', height=10).map(plt.scatter, "dim1", "dim2")
plt.title('TSNE - Amazon Review Dataset with (15k points, 30 Perplexity, 5k iterations)')
plt.grid()
plt.show()
```

```
plotTSNEModelForTFIDFW2V(sentVectors, final_df['Score'])
```

Converting to ndarray

Standardizing Array

Starting TSNE Model..

TSNE Model completed for TF-IDFW2V in 2054.895176887512



1.4 Conclusion / Final thoughts -

1. Each of the plots above for BOW, TFIDF, BowAvgw2V and TFIDFW2v methods are highly overlapped.
2. None of the plots gives a well separation between both positive/negative reviews.
3. We need to explore more using advance algorithms for better understanding or better differentiation for the classification.